

Exercise 1

1 Hypercube

- The labels in a d -dimensional hypercube use d bits. Fixing any k of these bits, show that the nodes whose labels differ in the remaining $d - k$ bit positions form a $(d - k)$ -dimensional sub-cube composed of 2^{d-k} nodes.
- One of the drawbacks of a hypercube-connected network is that different wires in the network are of different lengths (resulting in different data transfer times). It appears that 2D mesh networks with wraparound connections suffer from this drawback too. However, it is possible to fabricate a 2D wraparound mesh using wires of fixed length. Illustrate this layout by drawing such a 4×4 wraparound mesh.

2 Summation with OpenMP

In this exercise we want to sum up all natural numbers between 1 and 10^7 within a parallel for loop. Of course, there is no need to do this in a real application, because we can always explicitly compute the result without any iteration. However, this simple example let us test different approaches in handling the synchronized access on shared variables:

```
long long sumPar1(const int n) {  
    long long sum = 0;  
    #pragma omp parallel for default(none) shared(sum)  
    for (int i=1; i <= n; i++) {  
        sum += i;  
    }  
    return sum;  
}
```

- List at least four possibilities that OpenMP provides for making the given source code correct.
- Program the four different variants in summation.cpp, measure the performance, and check the results by comparing them with the explicit computed result.

3 Image Processing

Parallelizing sub-optimal serial codes often has undesirable effects of unreliable speedups and misleading runtimes. For this reason, we advocate optimizing serial performance of codes before attempting parallelization.

"imageprocessing.cpp" contains in the subroutine processSerial(...) an image processing program for edge detection in RGB images. For handling images we use the portable and open source library *FreeImage* (<http://freeimage.sourceforge.net/>).

- Compile (without any optimization flags), link and run the given program. Measure the runtime of this base variant (version 0). Since runtime measurements are always related to the performance of your test machine, write down the most important technical properties of your test machine.
- Study the compiler optimizations of your compiler and compile the same program with the best compiler settings. Compute the speedup of this version 1 to version 0.
- Now, have a closer look to the given program code. Copy the code in processSerial(...) to processSerialOpt(...) and try to optimize the new code as much as possible, but don't use parallelization. Use again the best suited compiler settings and compute the speedup of your optimized code compared to version 1.

Hints: The *FreeImage* methods getPixelColor(...) and setPixelColor(...) are very convenient but also very slow. Use instead getScanLine(...), but try to minimize the number of calls of this subroutine. A speedup to version 1 of ca. 6 should be possible.

- The optimized serial implementation (version 1) should be a good starting point for parallelization. Try to parallelize the outer most for-loop. You can use OpenMP or the parallelized for-loop in C++11. The expected speedup to version 1 should be almost the number of CPU cores in your machine, because there is no need for synchronization.